

IDEMPOTENCE

WHEN PERFORMING AN OPERATION AGAIN GIVES THE SAME RESULT

IDEMPOTENT

LOOK_AT_CAKE



LOOK_AT_CAKE



LOOK_AT_CAKE



LOOK_AT_CAKE



NOT IDEMPOTENT

EAT_SLICE_OF_CAKE

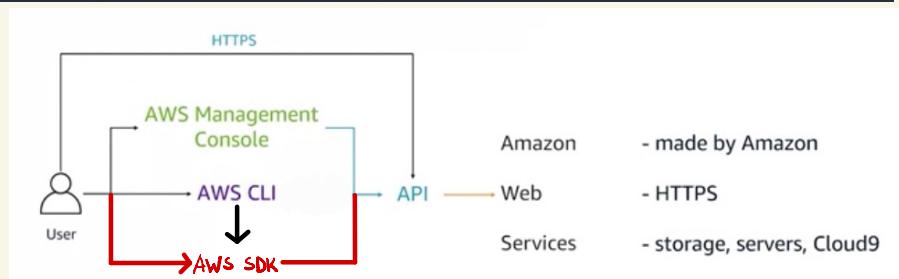
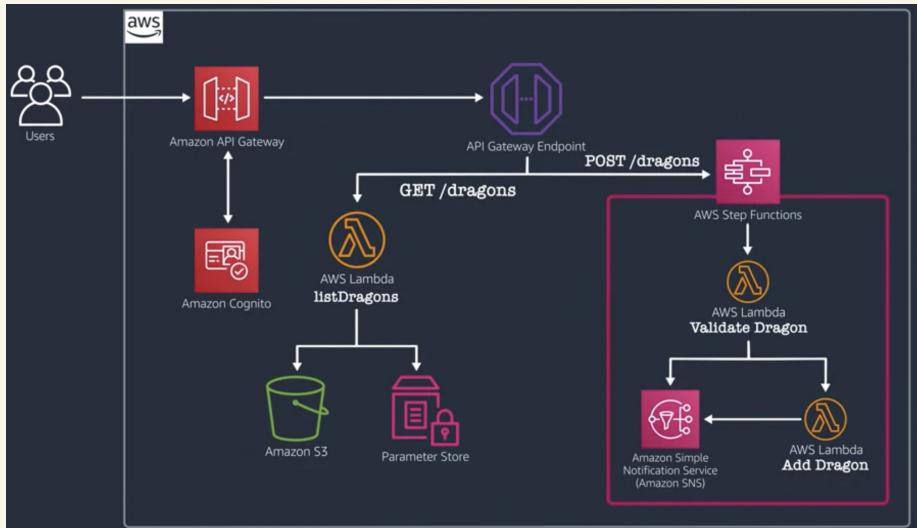
EAT_SLICE_OF_CAKE

EAT_SLICE_OF_CAKE

EAT_SLICE_OF_CAKE

EAT_SLICE_OF_CAKE





aws --region ca-central-1 s3 mb s3://modern-celapp

option command parameter sub command

→ Ways to interact with BOTO3

Client Interface

→ Low level

→ One-to-one map to AWS API

→ Responses come as python dictionaries

Resource Interface

→ High Level

→ Object-oriented

→ Wrapper for client

→ Exposes a subset of AWS API's

AWS

Client → import boto3

client = boto3.client('s3')

response = client.list_objects(Bucket='aws')

for content in response['Contents']:
 obj_dict = client.get_object(Bucket=

'aws', Key=content['Key'])

print(content['Key'], obj_dict['Last Modified'])

Resource → import boto3

Returns data in the form of object

s3 = boto3.resource('s3')

bucket = s3.Bucket('aws')

for obj in bucket.objects.all():
 print(obj.key, obj.last_modified)

access client from resource

s3_client = boto3.resource('s3').meta.client

Session - aws credentials

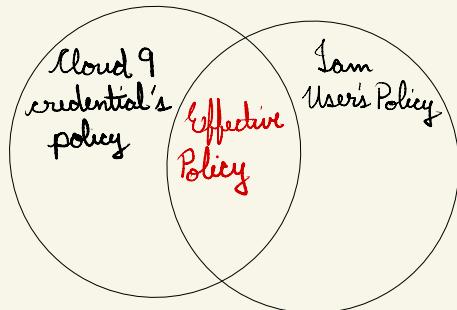
aws region we are submitting our API call

S3 select - SQL query read / filter a subset of data in an object

Data filtering done by S3

aws system parameter manager store: store / retrieve key-value pair

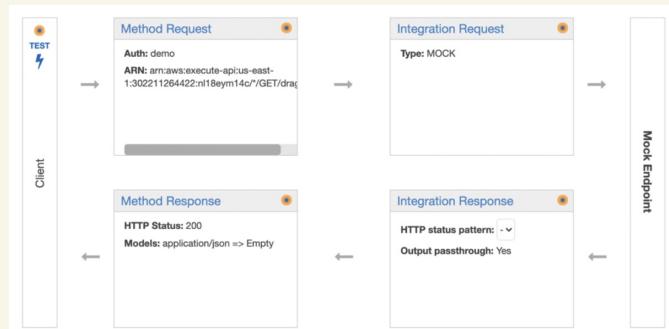
- Credentials passed into the client
- Credentials passed into the session
- Environmental Variable
- Credentials file
- IAM role
- cat ~/.aws/credentials } reset every 5 min



API gateway :-

- Amazon API enables you to create, publish, maintain, monitor, & secure your REST, HTTP & websocket API
- REST API ⇒ fast, stateless, standard, horizontal scalable, & dependable APIs
- Features ⇒ Canary deployment, Cloud Tail, CloudWatch, custom domain support, AWS service
- Methods :- client - facing interface by which calls can be made
- Edge-optimised API endpoints are best for geographical client.

- Regional API endpoints are intended for clients in the same region.
- Private API endpoints allows a client to securely access private API resources inside a VPC



API Gateway REST API → Validates requests/responses with model
 → Transforms requests/responses with mapping

Models :- Creating using JSON schema
 → Can be used to do request validation

Mapping :- → Backend expect a diff format than the client is sending. API gateway handles this
 → exceptional proxy integration
 → Written in VTL → supports XML
 → inject new parameters → conditional statement
 → hard code data → reference data in run time
 → map data in complex structure

Mock integration :- Enables your API to return a response for a request directly.

`https://[restapi-id].execute-api.[region].amazonaws.com/[stageName]`

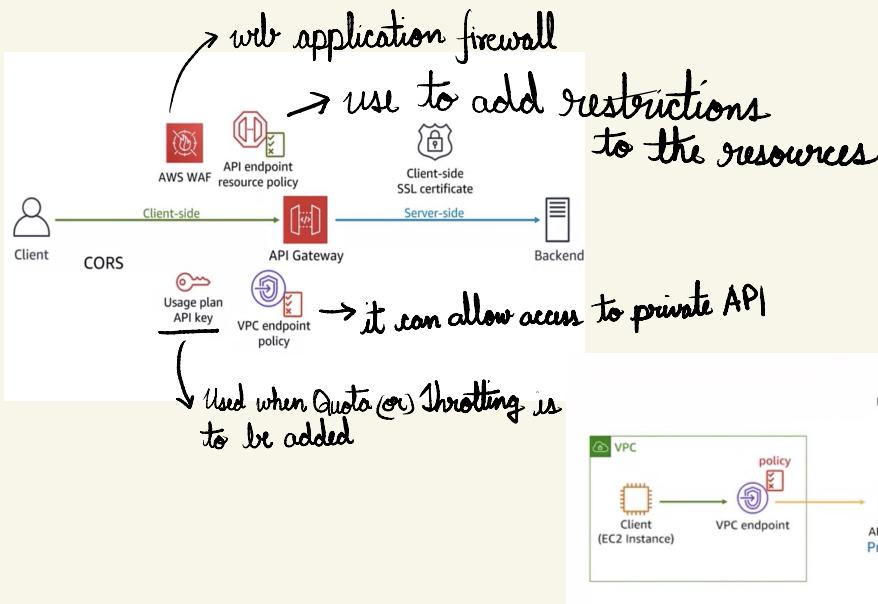
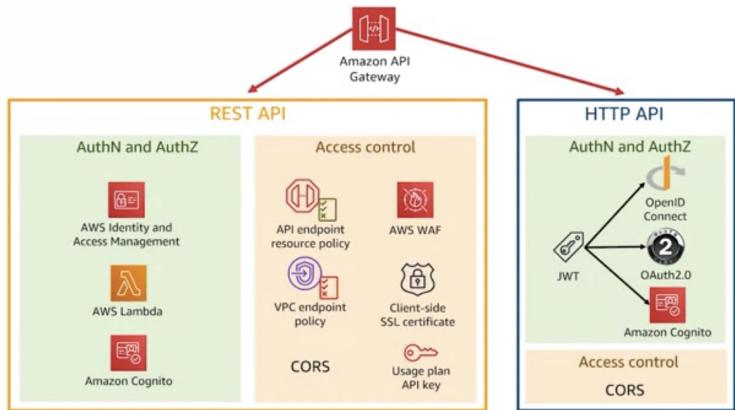
Stages Provide

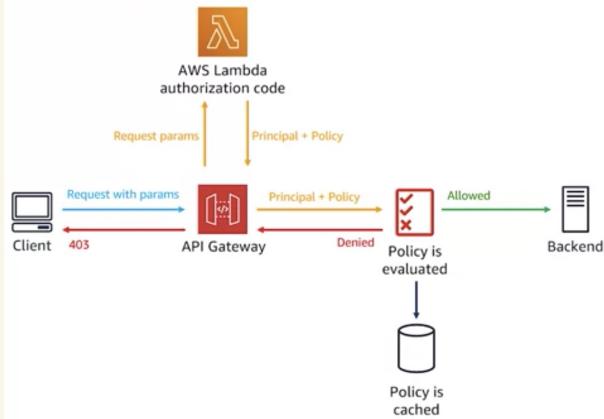
Method for deployment

API versioning

Performance management

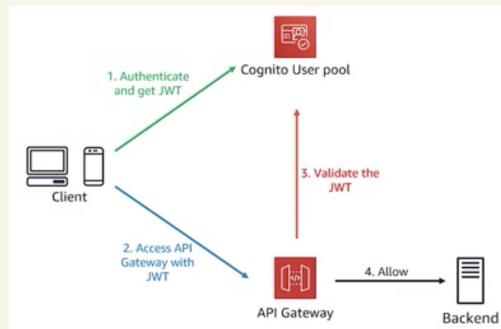
SDK generation





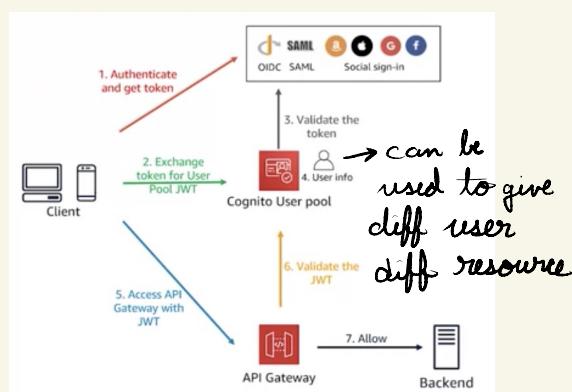
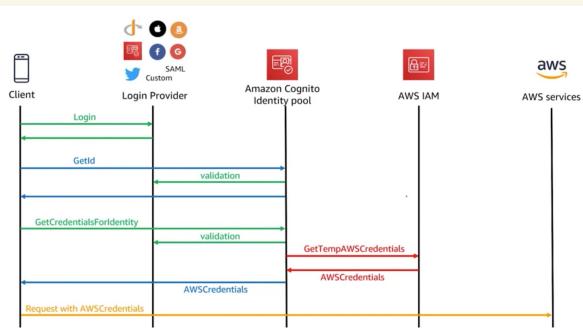
Lambda
→ can be customised

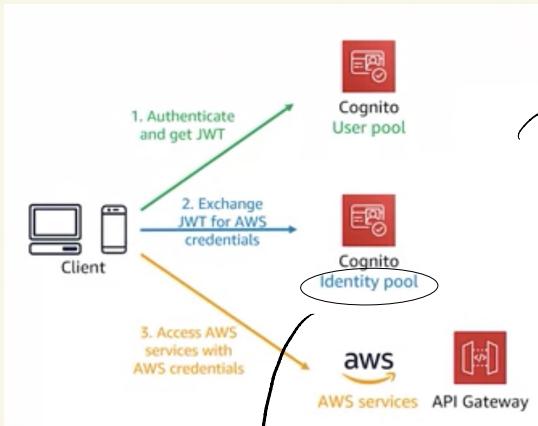
Loginto



- custom code can be run once someone signs up.
- allow backend to use data once sign up

→ When using a 3rd party for authentication





→ Federated identity giving basics IAM role to the not authorized user.

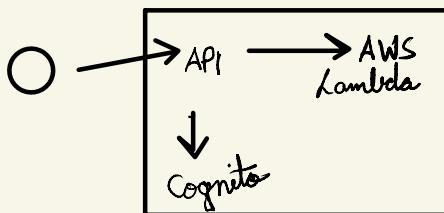
Aws Lambda

Event → JSON-formatted document containing data for a function to process.

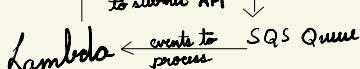
Concurrency → Number of requests a function is serving at any given time

Trigger → resource or configuration that invokes function

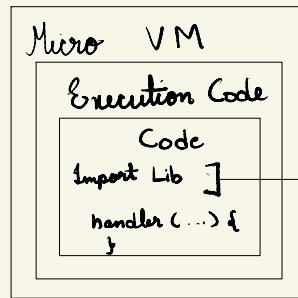
Push



Pull

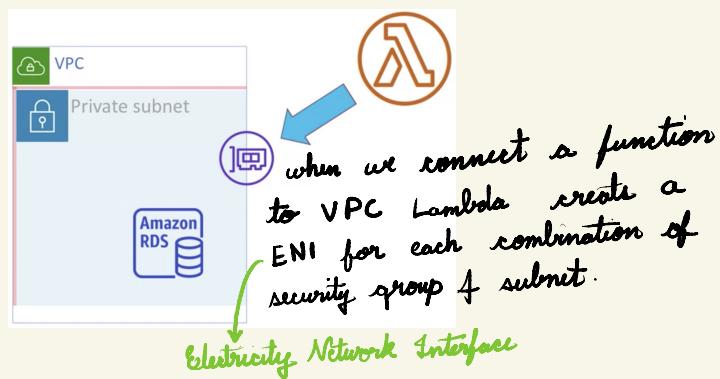


→ Provisioned
Concurrency

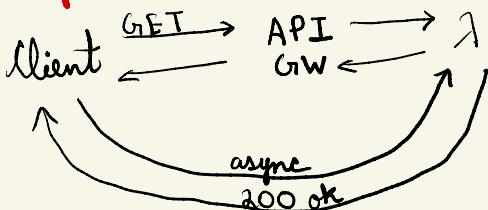


→ Bootstrap
executes everytime after cold
start

- Lambda endpoints only support secure connection over HTTPS
- AWS Lambda encrypts environment variables at rest
- Key config. allows to use an encryption key managed by AWS key management service.
- Encryption helper enables encryption of environment var. on the client side.



Synchronous + Asynchronous



Asynchronous Config

- Max. age of event
- Retry attempts → 1 or 2 only if it's idempotent
- Dead letter queue

can be applied multiple times without changing

(or) Lambda Destination

- asynchronous
- stream → pull model type of invocation with the trigger

Lambda function version

The system creates a new version each time a function is published.

- code, runtime, environmental variables, ARN
- we can use AWS management console ^{amazon} resource name
- (or) published version API
- every lambda version has a unique ARN
 - ARN for both function + version can't be changed once created
 - one for the function + one that points to specific version of the function.

Alias

- Pointer to a specific AWS Lambda function version accessed by its ARN.
- can be used when we have sources such as Amazon S3 event, which are used to invoke your lambda function.
- to create alias we use create alias API.

Boto3 + Lambda

```
import boto3
```

```
s3 = boto3.resource('s3', 'us-east-1').meta.client  
ssm = boto3.client('ssm', 'us-east-1')  
bucket_name = ssm.get_parameters()['Parameters'][0]['Value']  
file_name = " " (" ") " "
```

```
def listDragons(event, context):  
    expression = "select * from s3object s"
```

if 'queryStringParameters' in event:
 if 'query' in event['queryStringParameters']:
 query = event['queryStringParameters']['query']
 else:
 query = None

{ change expression

result = s3.select_object_content(
 Bucket = bucket_name
 Key = file_name
 ExpressionType = 'SQL')

Expression = expression,

InputSerialization = {'JSON': {'Type': 'Document'}}

OutputSerialization = {'JSON': {}}

records = " " ,

for event in result['Payload']:

if 'Records' in event:

records = event['Records'][0]['Payload']

decode('utf-8')

return {"statusCode": 200,

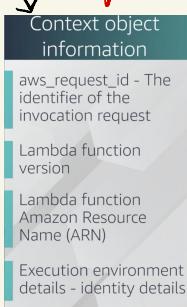
"body": json.dumps(records)}

```
→ sudo pip install --target ./t1 boto3
→ cd t1
→ zip -r9 ${OLDPWD}/ file.zip
→ zip -g file.zip l.py
→ aws lambda create-function --function-name <>
--runtime <> --role <> --handler <> --publish
--zip-file <>
→ aws lambda invoke --function --function-name <>
output.txt
→ aws lambda update-function-code --function-name <>
--zip-file <>
```

Python handler

```
def handler_name(event, context)
    # application logic
    return some_value
```

info coming in from
the request
(info on the object)



(info about the invocation
or the execution)

synchronous → JSON

asynchronous → the returned value
will be discarded

→ Boto 3 package is present in lambda but it must be present in zip form as it will update in future and can affect the perf.

Aws Step Function

- Helps to create server workflow
- Workflow in step function is called state machine
- State machine are composed of state
- Keep tracks of state machine & its execution.

Advantages

→ No server → auto scaling → parallel execution

It is built using amazon states language.

JSON-based, structured language used to define state machines & their workflow

State are individual ⇒ input → actions → output

Tasks - one unit of work done

Blue Point of Step Function

"Resource":
 "Catch": [
 {
 "ErrorEquals": "ErrorEquals", "Next": "Next", "ResultPath": "ResultPath"
 }
]

"Parameters": [
 {
 "Message": "Message", "PhoneNumber.\$": "PhoneNumber.\$"
 }
]

"Comment": "", "StartAt": "S1", "States": {
 "S1": {
 "Type": "Pass", "Result": "S1", "Next": "S2"
 }
 }

"Retry": [
 {
 "ErrorEquals": ["ErrorEquals"], "IntervalSeconds": "IntervalSeconds", "MaxAttempts": "MaxAttempts", "BackoffRate": "BackoffRate", "Start": "Start"
 }
]

"Task": [
 {
 "Task": "S1", "Choice": "IS+BOF", "Order": 1
 }
]

"Task": [
 {
 "Task": "S2", "Choice": "IS+BOF", "Order": 2
 }
]

"Task": [
 {
 "Task": "End", "Order": 3
 }
]

→ Path = string starting with \$

Types of state

- Types of state**

 - Task - do some real work [call other AWS service (integration)]
 - Choice - takes decision based on inputs
 - Parallel - when one state is not dependent on another for inputs ["Branches"]
 - Loop - more than one set of data [array]
→ combination of choice

- Map = we want to pass set of input concurrently [Man Concurrency:]
- Pass = takes input & pass it as output [testing]
 - if payload is not working at starting set it in middle to receive an output
 - we can use it while building state machine & then we can replace with appropriate type
- Wait → "Seconds"
- Succeed → ends the execution with a success
- Fail → " " " " " failure
 - [can specify the cause & error]
 - can be used in catch statement afterwards

AWS Step Function Service Integrations

```

    "Invoke Lambda function": {
        "Type": "Task",
        "Resource": "arn:aws:states:::lambda:invoke",
        "Parameters": {
            "FunctionName": "arn:aws:lambda:<REGION>:<ACCTNUMBER>:function:<FUNCTIONNAME>",
            "Payload": {
                "Input.$": "$"
            }
        },
        "Next": "NEXT_STATE"
    }
  
```

use to identify API action for
 the invoke lambda
 provided to the state

```
"Invoke Lambda State": {  
    "Type": "Task",  
    "Resource": "arn:aws:lambda:<REGION>:<ACCTNUMBER>:function:<FUNCTIONNAME>",  
    "Next": "NEXT_STATE"  
}
```

→ Amazon Dynamo DB [NoSQL]

```
"Put item into DynamoDB": {  
    "Type": "Task",  
    "Resource": "arn:aws:states:::dynamodb:putItem",  
    "Parameters": {  
        "TableName": "MyDynamoDBTable",  
        "Item": {  
            "DragonName": {  
                "S": "Atlas"  
            },  
            "Family": {  
                "S": "Red"  
            }  
        },  
        "Next": "NEXT_STATE"  
    }  
}
```

→ Amazon SNS

```
"Amazon SNS: Publish a message": {  
    "Type": "Task",  
    "Resource": "arn:aws:states:::sns:publish",  
    "Parameters": {  
        "Message": "Hello from Step Functions!",  
        "PhoneNumber": "+15555555555"  
    },  
    "Next": "NEXT_STATE"  
}
```

More step functions integrations

AWS Batch

Amazon ECS

Amazon SQS

AWS Glue

Amazon SageMaker

Amazon EMR

Step Function
AWS Lambda] Write Code
Internal resources]

Integration Pattern

You can call a service, and let step functions progress to the next state immediately after it gets an HTTP response. You can call a service and have step functions wait for job to complete or you can call a service with a task token and have step functions wait until that token is returned with a payload

⇒ API gateway → Lambda function → Step function → step machine
 Mapping uses VTL template.

console.aws.amazon.com/apigateway/home?region=us-east-1#/apis/it9pqbz5m7/resources/eqdk33/methods/POST

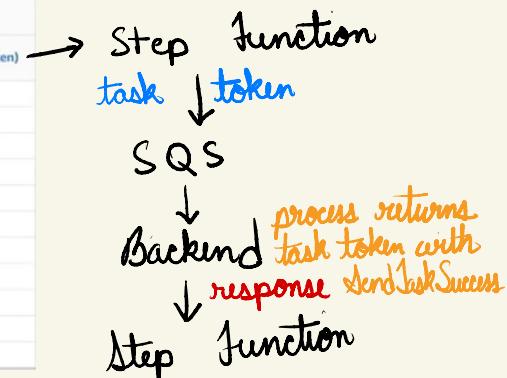
```

Generate template: [dropdown]
1 #set($data = $input.path('$'))
2
3 #set($input = " {"dragon_name_str": "$data.dragonName", "description_str": "$data.description", "family_str": "$data.family", "location_city_str": "$data.city", "location_country_str": "$data.country", "location_state_str": "$data.state", "location_neighborhood_str": "$data.neighborhood", "reportingPhoneNumber": "$data.reportingPhoneNumber", "confirmationRequired": $data.confirmationRequired}")
4
5
6 {
7     "input": "$Util.escapeJavaScript($input).replaceAll("\\\\", "")",
8     |"stateMachineArn": "arn:aws:states:us-east-1:302211264422:stateMachine:MyStateMachine"
9 }

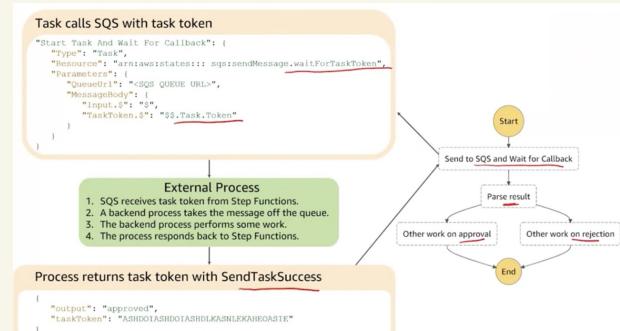
```

Supported Service Integrations

Service	Request Response	Run a Job (.sync)	Wait for Callback (.waitForTaskToken)
Lambda	✓		✓
AWS Batch	✓	✓	
DynamoDB	✓		
Amazon ECS/AWS Fargate	✓	✓	✓
Amazon SNS	✓		✓
Amazon SQS	✓		✓
AWS Glue	✓	✓	
Amazon SageMaker	✓	✓	
Amazon EMR	✓	✓	
AWS CodeBuild	✓	✓	
AWS Step Functions	✓	✓	✓

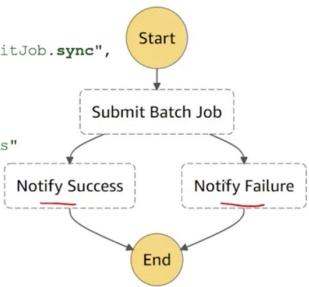


Step Function
 ↓
 SQS (Step Function writes the inputs are accepted by SQS)
 ↓
 Step Function proceeds to other task



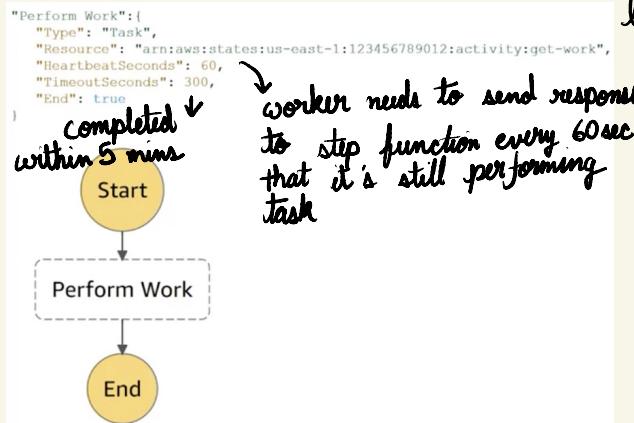
Run a job

```
"Submit Batch Job": {
    "Type": "Task",
    "Resource": "arn:aws:states:::batch:submitJob.sync",
    "Parameters": {
        "JobDefinition": "preprocessing",
        "JobName": "preprocessingBatchJob",
        "JobQueue": "primaryQueue",
        "Parameters.$": "$.batchjob.parameters"
    }
}
```



Step Function Activities

Any HTTPS protocol can be back end.



get activity task (will run infinitely)
 {ActivityArn: "..."}

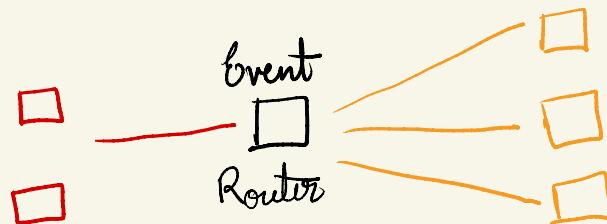
{ "input": " ", "taskToken: " " }
 Send Task Heartbeat
 { "taskToken": " " }
 Set Task Success (or) Failure
 { "error": "(or)" "output": " " }, taskToken: "

	Standard workflow	Express workflow
Execution start rate	Over 2,000 per second	Over 100,000 per second
State transition rate	Over 4,000 per second per account	Nearly unlimited
State transition persistence	Persisted to disk	In memory only
Execution semantics	Exactly-once workflow execution	At-least-once workflow execution
Pricing	Priced per state transition	Priced by the number of executions you run, their duration, and memory consumption
Maximum duration	365 days	5 minutes
Execution history	Stored in Step Functions and can be sent to Amazon CloudWatch Logs	Sent to Amazon CloudWatch Logs
Service integrations	Supports all service integrations and patterns	Supports all service integrations. Does not support Job-run (.sync) or Callback (.waitForTaskToken) patterns.

Event Driven architecture

↳ uses event to trigger + communicate between decoupled services.

change in a state



→ Producer + consumer services are decoupled

→ Less **orchestration** than step function

↳ the planning of the elements of a situation to produce a desired effect

→ **SQS** → Simple Queue Services

↳ we can send, store + receive messages
14 days
Standard SQS [Best effort ordering]
→ idempotent
SQS FIFO

→ **SNS** → Producer can send message via the
HTTP protocol, fan out message to
large number of consumer end point for n process

→ high throughput
→ lower latency
event

Event Bridge

↳ A service bus that makes it easy to connect producers & consumers together using data from the application
→ Setup routing rules
→ Schema Registry → It stores event structure or schema in a shared central location that consumer can query.

AWS X-RAY

Segment → Compute resources running the application logic
→ resource's name, details about the request, work done
→ can breakdown the data about the work done into subsegments.
 ↳ timing information & downstream calls
→ In Amazon DB, X-Ray uses subsegments to generate inferred segments & downstream nodes on service map.

Service graph → data sent from application (JSON)

Trace - tracking of the path of a request through the system collecting the segments generated by the request.

→ X-Ray provides a user-centric model.

- API gateway - X-Ray
- AWS Lambda - X-Ray Daemon
- AWS CloudTrail - X-Ray
- AWS CloudWatch
 - EC2
 - Elastic Beanstalk
 - SNS → SQS

Tracing →
passive [has been enabled on upstream]
active [auto based on sampling algo]

X-Ray : Python

X-Ray

Trade data

Incoming HTTP requests

Downstream calls made using the AWS SDK

HTTP clients

SQL database connectors

SDK for python → creating & sending trace data to X-Ray daemon.
(run on UDP port 2000)

Annotations allow you to associate key-value pairs with a trace that can then be filtered in the X-ray console.

Metadata information can be used to store data that can provide further information

Subsegments

Allows you to provide more granular detail about a segment, including custom information

Boto3 + Lambda

[X-Ray]

allow us to record all the downstream calls
 import boto3 to other AWS services
 from aws-xray-sdk.core import patch_all
 patch_all()

```
s3 = boto3.resource('s3', 'us-east-1').meta.client
```

```
ssm = boto3.client('ssm', 'us-east-1')
```

```
bucket_name = ssm.get_parameters()['Parameters'][0]['Value']
```

```
file_name = " " (" ") " "
```

def listDragons(event, context):

expression = "select * from s3object \$"

if 'queryString Parameter' in event['querystring']
 not None:

{ change expression

result = s3.select_object_content(

Bucket = bucket_name

Key = file_name

ExpressionType = 'SQL'

Expression = expression,

InputSerialization = {'JSON': {'Type': 'Document'}}

Output = {"": {"JSON": {"Type": "Document"}}}

records = ````

for event in result['Payload']:

if 'Records' in event:

records = event['Records'][0]['Payload'].

decode('utf-8')

return {"statusCode": 200,

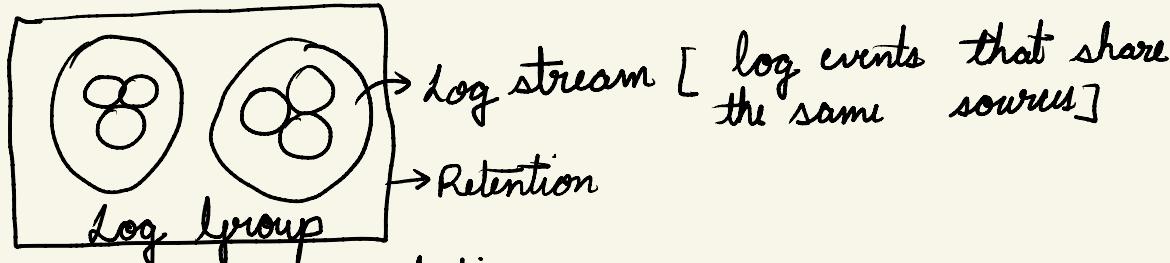
"body": json.dumps(records)}

calling
select object
API

→ sudo pip install aws-ray-sdk
→ aws lambda update-function-configuration
-- function-name list-dragons --tracing-config
Mock = Active

Log event - This is a record of some activity recorded by the application.

→ Timestamp of the event + in the raw event message



[share the same retention monitoring or access control settings]

Logs are sent through HTTPs request

API gateway
Step Function
Lambda } Logs → Cloudwatch Logs

API gateway API
Login

give permissions

AWS Identity & Access Management Services

[IM role with IM policy]

execution logging → log data includes errors, [request + resp] payloads, data used by lambda authorizers, whether API keys were required, " usage plans are enabled

access logging → REST API

→ who has accessed & how he has accessed

Express workflow - default

Standard workflow - need to enable manually

Edge → [less physical space]

↳ CloudFront [caching service]

Invocation types

Viewer request to Amazon CloudFront

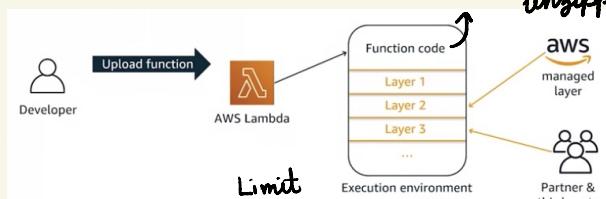
Viewer request to origin

Origin response to Amazon CloudFront

Origin response to viewer

→ Lambda

[128 - 3008 MB]



can't exceed 250MB unzipped

Create a layer [publish layer version command]

→ name

→ description

→ zip archive

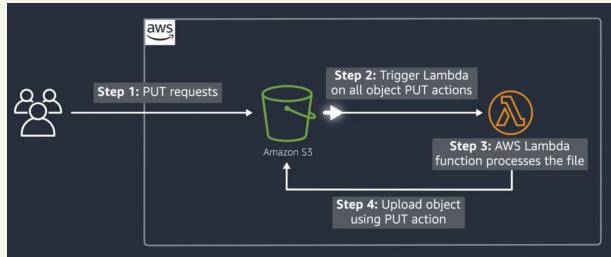
→ Lambda Best Practices

→ Min. package size

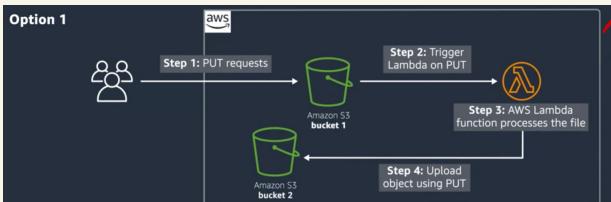
[75GB]

→ Min. the complexity of your dependencies

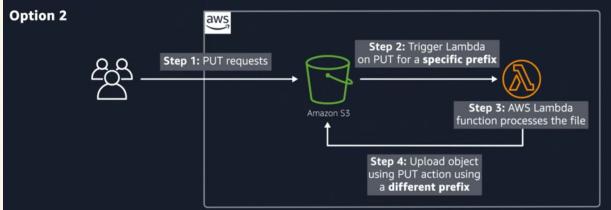
→ Avoid using recursive call



∞ loop



Sol.



Amazon API Gateway HTTP APIs
Lambda proxy
HTTP endpoint proxy
OpenID Connect
OAuth 2.0 for authorization
CORS support
Private integration support

→ set concurrent execution limit to 0 & update the code

→ Take adv. of execution content reuse

don't store user data
(or) sensitive data in execution content

→ Load test your lambda to determine optimal timeout value.